

VAIBHAV GOVILKAR

Sunnyvale, CA · (773) 844-3097 · vaibhavgovilkar88@gmail.com · linkedin.com/in/vaibhavgovilkar

SUMMARY

Staff Software Engineer and Tech Lead with 12+ years of experience driving technical strategy, building high-impact platform teams, and delivering developer-facing products at scale. Currently leading the AI model deployment platform at Meta — serving as both TL and de facto PM for an enterprise-scale AI model deployment platform — defining roadmap, metrics, and customer engagement strategy. Drove unassisted deployment rate from 37% to 66% and deployment plan approval rate from 10% to 33% through strategic bets, cross-functional alignment, and tight customer feedback loops. Currently pioneering AI-native engineering workflows — designing autonomous agent systems, ai-native user experiences, and orchestration architectures that enable engineers to dramatically improve throughput. Track record of turning ambiguous problem spaces into structured roadmaps, mentoring engineers, and influencing org-wide direction.

EXPERIENCE

Meta | Menlo Park, CA Jan 2017 – Present

9+ year tenure spanning AI infrastructure, VR/AR, and social platforms. Joined as a founding engineer on Horizon Worlds, promoted twice from mid-level engineer to Staff Engineer / Tech Lead. Grew from individual contributor to owning team strategy, roadmap, and customer engagement across orgs. Conducted 108+ technical interviews and supported hiring decisions and headcount planning.

Staff Engineer & Tech Lead — AI Inference Aug 2024 – Present

AI Model Deployment Platform

- Own end-to-end strategy, roadmap, and execution for Meta's AI model deployment platform — functioning as both Tech Lead and de facto Product Manager in the absence of a dedicated PM, spanning product direction, UX scoping, stakeholder alignment, and engineering delivery.
- Drove key platform outcomes: unassisted deployment rate 37%→66%, deployment plan approval rate 10%→33%, CSAT 48%→57% — by identifying high-leverage investments, defining success metrics with data science, and ruthlessly prioritizing the roadmap against customer pain points. These self-serve improvements enabled the platform to absorb a 2x YOY increase in model launch volume without proportional headcount growth.
- Established and maintained tight feedback loops with internal customers (ML engineers shipping models across Instagram, Facebook & Threads) through regular syncs, stakeholder surveys, deployment retrospectives, and escalation management — translating user friction into prioritized workstreams.
- Led a cross-functional task force that achieved a 14% improvement in unassisted deployments — aligning partners on process, commitments, and exit criteria to systematically eliminate automation blockers.
- Defined deployment SLOs and evolved key efficiency metrics from self-reported to signal-based measurement — collaborating with data science and data engineering to eliminate reporting bias and establish shared goals with leadership and capacity teams.
- Set strategic direction for capacity & quota automation, model-readiness gating, and configuration management — authoring RFCs, aligning with leadership, and shaping the long-term platform vision.
- Led the design and implementation of AI-powered experiences embedded in the deployment platform — including a domain agent exposing launch operations via MCP tools, a context-aware agent side panel in the launch wizard, and a 'Resolve with Agent' pattern that routes users directly to AI agents for automated issue resolution.
- Architected a persistent AI engineering orchestration system operating as a continuous force multiplier — including cron-based automated task resolution with CI monitoring and self-healing diffs, parallel AI coding agents on isolated checkpoints, real-time chat-to-implementation pipelines across 50+ channels, and daily communication synthesis into structured action items.
- Developed an AI-assisted development methodology for platform reliability — authored launch validation suites and reliability strategies through a specification-driven loop where AI agents generate implementations, tests, and schema changes, demonstrating the architect-plus-agent-fleet operating model.
- Drove production excellence culture: owned/supported resolution of 17 production incidents, supported 115 model deployments, instituted incident severity guides, established post-incident review processes, and created peak-season policies including code freezes and on-call surge groups.
- Managed half-year planning cycles end-to-end — brainstorming sessions, customer leadership alignment surveys, workstream scoping for resourcing, and quarterly lookbacks with roadmap pivots communicated to stakeholders.
- Mentored and grew engineers through regular 1:1s, technical and product direction, and progressive delegation; represented the team in cross-org leadership forums, ran knowledge-sharing summits, and navigated difficult customer escalations.

Staff Engineer & Tech Lead — Reality Labs Jan 2022 – Aug 2024

VR World Building Tools & Creator Platform

- Tech led a major engine upgrade rollout end-to-end — owned rollout strategy, communications, and engineering; pivoted to war-room leadership when needed, aligning design and partner teams on critical UX changes, costing solutions, and driving rapid iteration loops to unblock each release phase through to general availability.
- Unblocked stalled product launches by identifying creative product pathways — designed new content types and self-serve tooling to ship a 3D model import pipeline to external creators (152 projects published), and managed difficult partner communications to drive legacy platform migration to completion.
- Led technical direction and roadmapping for world streaming architecture (~3x load-time improvement), multi-material rendering support, and a generative AI mesh creation prototype — maintaining active feedback loops with creators to inform spec changes and rollout adjustments.
- Drove platform reliability as largest contributor to error-rate reduction (25%→~0%) on high-volume operations, hitting platform reliability goals; led a cross-functional effort to delete 187K lines of legacy code (~3MB application binary reduction).
- Led developer productivity strategy — evaluated and documented acceleration options, ran a feasibility study on a lightweight editor approach, and conducted an org-wide survey to estimate impact; drove annual planning, roadmapping, and GenAI resourcing.
- Set the team testing strategy and served on the org's build health and stability working group — a TL-staffed on-call rotation enforcing engineering standards and product stability; primary engineering point person and technical consultant across the creator platform.

Software Engineer (Founding Engineer) — Horizon Worlds Jan 2017 – Dec 2021

VR Social Platform

- One of 10 founding engineers on Horizon Worlds — helped build the team and product from inception through rapid growth to a multi-hundred person organization, shipping core platform features in PHP/Hack, C#, Unity3D, and React and establishing early engineering patterns adopted as the team scaled.

Software Engineer | Orbitz Worldwide (Expedia) Jan 2014 – Jan 2017

Chicago, IL

- Architected microservices decomposition of monolithic SEO platform into a service-oriented stack with automated Docker deployment pipelines.
- Built real-time flight search pricing pipeline (ElasticSearch) and recommendation engine for cheapest historic air prices on SEO landing pages.
- Led site-awareness automation overhaul — continuous sitemap generation, SEO monitoring via Google Search Console API, and automated landing page updates.

Software Engineer | Nokia Corporation Jun 2013 – Dec 2013

Chicago, IL

- Developed client-server features, analytics, and UI enhancements for Nokia Xpress mobile browser.

Research Assistant | University of Illinois at Chicago May 2012 – May 2013

Electronic Visualization Lab

- Built rail-track simulation visualization for the CAVE2 VR environment using C++, OpenGL, GLSL shaders, and OmegaLib.

EDUCATION

M.S. Computer Science — University of Illinois at Chicago | GPA 3.62 2011 – 2013

MBA + B.Tech — NMIMS University, Mumbai 2006 – 2011

TECHNICAL SKILLS

Languages: Python, JavaScript/TypeScript, C++, C#, PHP/Hack, Java, Bash

Frameworks & Tools: React, Unity3D, GraphQL, Thrift, Docker, ElasticSearch, Git, OpenGL/GLSL

Domains: AI/ML Infrastructure, Model Deployment & Serving, Agentic Systems & AI-Native Tooling, Platform & Developer Tooling, VR/3D Graphics